



auto_offset_earliest

Simon Bärle, Nicolas Neudeck



Tech-Stack





Task 1

- `org.apache.kafka.clients.producer.KafkaProducer`
- `kotlinx.serialization.json`
- Create topic with Conduktor (UI) or broker bash (“events”)
- Read JSON file line by line:
 - Parse Line and extract ID field from event (use as Event ID)
 - Convert Line String to `ByteArray` and send event to Kafka
 - `producer.serializer = ByteArraySerializer`

```
fun produce(topicName:String, value: String){  
    val key = json.parse(Data.serializer(), value).id  
    val producerRecord = ProducerRecord(topicName, key, value.toByteArray())  
    kafkaProducer.send(producerRecord)  
}
```



Task 2

- Added ksqlDB and ksqlDB-CLI to setup
- Streams
 - meetup_events_stream
 - meetup_events_stream_de
 - meetup_events_stream_de_munich

```
CREATE STREAM meetup_events_stream_de AS SELECT name, "GROUP"->city  
AS city, "GROUP"->country AS country FROM meetup_events_stream  
WHERE "GROUP"->country = 'de';
```

```
CREATE STREAM meetup_events_stream_de_munich AS SELECT name, city, country  
FROM meetup_events_stream_de WHERE city = 'Munich' OR city = 'München';
```



Latency

- `org.apache.kafka.clients.producer.KafkaConsumer`
- Kafka automatically adds creation timestamp on event (per topic)
- Two concurrent consumers (using Coroutines)
 - Consumer 1 (“events”)
 - Consumer 2 (“meetup_events_stream_de_munich”)
 - Consumers store `event_id` and `event_timestamp` in map
- Compare timestamps for same `event_id` to get processing time of kafka/ksqlDB
- Result: Processing time of Kafka is very low (<1 ms)
- Potential error: Copied creation timestamp between different queues



Throughput & Kafka Metrics

- Initial Idea: Collect built-in metrics from Kafka Producer/Consumer
 - Multiple, inconsistent results per metric (no reliable data source!)
- Own Measurement in code
 - Measure execution time and event file size
- Built-in metrics from Conduktor
 - Interesting Outcome: Docker on Mac has substantially higher networking overload (way slower than on Linux)
- Results: produced full events.json (~350MB) in ~18 seconds to Kafka (~19,4 MB/s)
- Cf. Kafka Performance Evaluation (better hardware): ~78MB/s

Task 3 - Hands On

